


```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from google.colab import drive
import os

drive.mount("/content/drive", force_remount=True)

# Make sure to replace 'MyDrive/path/to/your/file.txt' with the actual path
file_path = '/content/drive/MyDrive/Colab Notebooks/MSDS 422/Module 1/bangladesh_data.csv'

```

 Mounted at /content/drive

## ✓ 1. EDA on Bangladesh's Socio-economic data

```

# Load dataset
data = pd.read_csv(file_path)
data.columns = data.columns.str.strip() # Clean column names

# Display basic information
print("Dataset Shape:", data.shape)

print("\nBasic Info:")
data.info()

print("\nSummary Statistics:")
print(data.describe())

print("\nFirst Few Rows:")
print(data.head())

print("\nMissing Values:")
print(data.isnull().sum())

```

 [Show hidden output](#)

```

# Columns of interest for analysis
columns_of_interest = [
    "Year",
    "Poverty Rate (National)",
    "Rural Poverty Rate",
    "Urban Poverty Rate",
    "Literacy Rate(%)",
    "GDP",
    "Education Spending (% of GDP)",
    "Healthcare Spending Per Capita (US$)",
    "Population",
    "Urban Population % of Total",
    "Unemployment Rate (%)",
    "Inflation Rate (%)",
    "Net Migration Rate",
    "GDP Per Capita",
    "Export Growth(%GDP)",
    "Import Growth(%GDP)",
    "Death Rate",
    "Birth Rate",
    "Infant Mortality Rate",
    "Life Expectancy Growth Rate (%)",
    "Annual % Crime Rate Change",
    "green house gas emissions Annual % Change",
    "Fossil Fuel consumption % of Total Energy Use",
    "Maternal Mortality Rate Per 100K Live Births",
    "Clean Water Access % of Population",
    "Electricity Access % of Population",
    "Suicide Rate"
]

len(columns_of_interest)

```

 27

```
# calculate missing value % for each feature
missing_percentage = (data.isnull().sum() / len(data)) * 100
print(missing_percentage.sort_values())
```

```
Year                                0.000000
Infant Mortality Rate               0.000000
Growth Rate.2                       0.000000
Unemployment Rate (%)              0.000000
Annual Change (%)                  0.000000
...
tourism exports                     63.636364
Tourism spending                    63.636364
Annual Change.4                     75.000000
Coal consumption % of Electricity from Coal  75.000000
Number of Private Vehicles          100.000000
Length: 118, dtype: float64
```

## ✓ 1.1 Preping Bangladesh's Socio-economic data and Feature Engineering data for further analysis

```
# Check if columns exist in the dataset
missing_columns = [col for col in columns_of_interest if col not in data.columns]
if missing_columns:
    print(f"\nColumns not found in the dataset: {missing_columns}")
else:
    # Subset and copy relevant data
    data_subset = data[columns_of_interest].copy()
    print("\nData Subset:")
    print(data_subset.head())

    # Clean and preprocess data
    # Handle percentage columns
    percentage_columns = ["Poverty Rate (National)", "Rural Poverty Rate", "Urban Poverty Rate", "Literacy Rate(%)", "Education
                          "Unemployment Rate (%)",
                          "Inflation Rate (%)",
                          "Net Migration Rate",
                          "Export Growth(%GDP)",
                          "Import Growth(%GDP)",
                          "Death Rate",
                          "Birth Rate",
                          "Infant Mortality Rate",
                          "Life Expectancy Growth Rate (%)",
                          "Annual % Crime Rate Change",
                          "green house gas emissions Annual % Change",
                          "Fossil Fuel consumption % of Total Energy Use",
                          "Maternal Mortality Rate Per 100K Live Births",
                          "Clean Water Access % of Population",
                          "Electricity Access % of Population",
                          "Suicide Rate"]
    for col in percentage_columns:
        if data_subset[col].dtype == 'object':
            data_subset[col] = data_subset[col].str.rstrip('%').astype('float') / 100.0

    # Clean GDP column
    if data_subset["GDP"].dtype == 'object':
        data_subset["GDP"] = data_subset["GDP"].replace(['\$',B'], '', regex=True).astype('float')
    if data_subset["GDP Per Capita"].dtype == 'object':
        data_subset["GDP Per Capita"] = data_subset["GDP Per Capita"].replace(['\$',B'], '', regex=True).astype('float')

    # Clean Healthcare Spending column
    if data_subset["Healthcare Spending Per Capita (US$)"].dtype == 'object':
        data_subset["Healthcare Spending Per Capita (US$)"] = data_subset["Healthcare Spending Per Capita (US$)"].replace(['\$',
                                                                                                     B'])

    # Clean Population column
    if data_subset["Population"].dtype == 'object':
        data_subset["Population"] = data_subset["Population"].replace('[,]', '', regex=True).astype('float')

    # Rename columns for consistency
    data_subset.columns = [
        "Year", "National Poverty Rate", "Rural Poverty Rate", "Urban Poverty Rate",
        "Literacy Rate", "GDP", "Education Spending", "Healthcare Spending",
        "Population", "Urban Population (%)",
        "Unemployment Rate",
        "Inflation Rate",
```

```

    "Net Migration Rate",
    "GDP Per Capita",
    "Export Growth",
    "Import Growth",
    "Death Rate",
    "Birth Rate",
    "Infant Mortality Rate",
    "Life Expectancy Growth Rate",
    "Annual Crime Rate Change",
    "Greenhouse Gas Emissions Change",
    "Fossil Fuel Consumption",
    "Maternal Mortality Rate",
    "Clean Water Access",
    "Electricity Access",
    "Suicide Rate"
]

# Impute missing values
imputer = SimpleImputer(strategy="mean")
data_subset.iloc[:, 1:] = imputer.fit_transform(data_subset.iloc[:, 1:])

# Feature Scaling
scaler = StandardScaler()
scaled_data = scaler.fit_transform(data_subset.iloc[:, 1:])
scaled_data = pd.DataFrame(scaled_data, columns=data_subset.columns[1:])
data_subset = pd.concat([data_subset["Year"], scaled_data], axis=1)

# Feature Engineering
data_subset["Urbanization × GDP Growth"] = data_subset["Urban Population (%)"] * data_subset["GDP"]
data_subset["3-Year GDP Rolling Avg"] = data_subset["GDP"].rolling(window=3).mean()
data_subset["3-Year Literacy Rolling Avg"] = data_subset["Literacy Rate"].rolling(window=3).mean()

```

 [Show hidden output](#)

```

# do the same data processing but keep original values without scaling

# Check if columns exist in the dataset
missing_columns = [col for col in columns_of_interest if col not in data.columns]
if missing_columns:
    print(f"\nColumns not found in the dataset: {missing_columns}")
else:
    # Subset and copy relevant data
    data_original = data[columns_of_interest].copy()
    print("\nData Subset:")
    print(data_original.head())

# Clean and preprocess data
# Handle percentage columns
percentage_columns = ["Poverty Rate (National)", "Rural Poverty Rate", "Urban Poverty Rate", "Literacy Rate(%)", "Education
    "Unemployment Rate (%)",
    "Inflation Rate (%)",
    "Net Migration Rate",
    "Export Growth(%GDP)",
    "Import Growth(%GDP)",
    "Death Rate",
    "Birth Rate",
    "Infant Mortality Rate",
    "Life Expectancy Growth Rate (%)",
    "Annual % Crime Rate Change",
    "green house gas emissions Annual % Change",
    "Fossil Fuel consumption % of Total Energy Use",
    "Maternal Mortality Rate Per 100K Live Births",
    "Clean Water Access % of Population",
    "Electricity Access % of Population",
    "Suicide Rate"]
for col in percentage_columns:
    if data_original[col].dtype == 'object':
        data_original[col] = data_original[col].str.rstrip('%').astype('float') / 100.0

# Clean GDP column
if data_original["GDP"].dtype == 'object':
    data_original["GDP"] = data_original["GDP"].replace(['\$', 'B'], '', regex=True).astype('float')
if data_original["GDP Per Capita"].dtype == 'object':
    data_original["GDP Per Capita"] = data_original["GDP Per Capita"].replace(['\$', 'B'], '', regex=True).astype('float')

# Clean Healthcare Spending column

```

```

if data_original["Healthcare Spending Per Capita (US$)"].dtype == 'object':
    data_original["Healthcare Spending Per Capita (US$)"] = data_original["Healthcare Spending Per Capita (US$)"].replace(

# Clean Population column
if data_original["Population"].dtype == 'object':
    data_original["Population"] = data_original["Population"].replace('[,]', '', regex=True).astype('float')

# Rename columns for consistency
data_original.columns = [
    "Year", "National Poverty Rate", "Rural Poverty Rate", "Urban Poverty Rate",
    "Literacy Rate", "GDP", "Education Spending", "Healthcare Spending",
    "Population", "Urban Population (%)",
    "Unemployment Rate",
    "Inflation Rate",
    "Net Migration Rate",
    "GDP Per Capita",
    "Export Growth",
    "Import Growth",
    "Death Rate",
    "Birth Rate",
    "Infant Mortality Rate",
    "Life Expectancy Growth Rate",
    "Annual Crime Rate Change",
    "Greenhouse Gas Emissions Change",
    "Fossil Fuel Consumption",
    "Maternal Mortality Rate",
    "Clean Water Access",
    "Electricity Access",
    "Suicide Rate"
]

# Impute missing values
imputer = SimpleImputer(strategy="mean")
data_original.iloc[:, 1:] = imputer.fit_transform(data_original.iloc[:, 1:])

# Feature Engineering
data_original["Urbanization x GDP Growth"] = data_original["Urban Population (%)"] * data_original["GDP"]
data_original["3-Year GDP Rolling Avg"] = data_original["GDP"].rolling(window=3).mean()
data_original["3-Year Literacy Rolling Avg"] = data_original["Literacy Rate"].rolling(window=3).mean()

```

 [Show hidden output](#)

```

print(data_subset.shape)
print(data_original.shape)

```


 [Show hidden output](#)

```

# create a subset of data with columns_of_interest
data_original_2 = data[columns_of_interest]

# calculate missing value % for each feature
missing_percentage = (data_original_2.isnull().sum() / len(data_original_2)) * 100
print(missing_percentage)

```

 Year 0.000000  
Poverty Rate (National) 0.000000  
Rural Poverty Rate 0.000000  
Urban Poverty Rate 0.000000  
Literacy Rate(%) 0.000000  
GDP 0.000000  
Education Spending (% of GDP) 18.181818  
Healthcare Spending Per Capita (US\$) 50.000000  
Population 0.000000  
Urban Population % of Total 0.000000  
Unemployment Rate (%) 0.000000  
Inflation Rate (%) 0.000000  
Net Migration Rate 2.272727  
GDP Per Capita 0.000000  
Export Growth(%GDP) 0.000000  
Import Growth(%GDP) 0.000000  
Death Rate 0.000000  
Birth Rate 0.000000  
Infant Mortality Rate 0.000000  
Life Expectancy Growth Rate (%) 0.000000  
Annual % Crime Rate Change 56.818182  
green house gas emissions Annual % Change 29.545455  
Fossil Fuel consumption % of Total Energy Use 20.454545  
Maternal Mortality Rate Per 100K Live Births 52.272727

Clean Water Access % of Population	47.727273
Electricity Access % of Population	27.272727
Suicide Rate	54.545455

dtype: float64

```
# calculate missing value % for each feature from imputed dataset
missing_percentage = (data_original.isnull().sum() / len(data_original)) * 100
print(missing_percentage)
```

```
Year 0.000000
National Poverty Rate 0.000000
Rural Poverty Rate 0.000000
Urban Poverty Rate 0.000000
Literacy Rate 0.000000
GDP 0.000000
Education Spending 0.000000
Healthcare Spending 0.000000
Population 0.000000
Urban Population (%) 0.000000
Unemployment Rate 0.000000
Inflation Rate 0.000000
Net Migration Rate 0.000000
GDP Per Capita 0.000000
Export Growth 0.000000
Import Growth 0.000000
Death Rate 0.000000
Birth Rate 0.000000
Infant Mortality Rate 0.000000
Life Expectancy Growth Rate 0.000000
Annual Crime Rate Change 0.000000
Greenhouse Gas Emissions Change 0.000000
Fossil Fuel Consumption 0.000000
Maternal Mortality Rate 0.000000
Clean Water Access 0.000000
Electricity Access 0.000000
Suicide Rate 0.000000
Urbanization × GDP Growth 0.000000
3-Year GDP Rolling Avg 4.545455
3-Year Literacy Rolling Avg 4.545455
dtype: float64
```

```
data_original['Healthcare Spending'].value_counts()
```




Healthcare Spending	count
25.545455	22
12.000000	2
9.000000	2
27.000000	1
51.000000	1
48.000000	1
45.000000	1
42.000000	1
40.000000	1
34.000000	1
30.000000	1
23.000000	1
24.000000	1
21.000000	1
18.000000	1
16.000000	1
14.000000	1
11.000000	1
10.000000	1
8.000000	1
58.000000	1


dtype: int64

## 1.2 Correlation Analysis on Features for data relations

```
# Exploratory Data Analysis
# Correlation Matrix
correlation_matrix = data_subset.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=False, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix (Including Created Features)")
plt.show()
```



```
# check correlation with National Poverty Rate
correlation_matrix["National Poverty Rate"].sort_values(ascending=False)
```





### National Poverty Rate

National Poverty Rate	1.000000
Rural Poverty Rate	0.999709
Urban Poverty Rate	0.998054
Infant Mortality Rate	0.992928
Death Rate	0.947170
Birth Rate	0.917431
Life Expectancy Growth Rate	0.775513
Net Migration Rate	0.535455
Maternal Mortality Rate	0.302007
Suicide Rate	0.276712
Annual Crime Rate Change	0.071631
Greenhouse Gas Emissions Change	-0.102458
Clean Water Access	-0.275660
Education Spending	-0.301818
Healthcare Spending	-0.311259
Inflation Rate	-0.373114
Urbanization × GDP Growth	-0.396160
Electricity Access	-0.574456
Import Growth	-0.652610
Fossil Fuel Consumption	-0.740503
Export Growth	-0.799517
GDP	-0.814500
GDP Per Capita	-0.818170
3-Year GDP Rolling Avg	-0.819651
Unemployment Rate	-0.907094
Literacy Rate	-0.968302
3-Year Literacy Rolling Avg	-0.973785
Urban Population (%)	-0.979221
Year	-0.997426
Population	-0.998727

dtype: float64

### 1.3 Feature trend graph analysis



```
# Trend Analysis for original data
import numpy as np
from matplotlib.ticker import MaxNLocator

plt.figure(figsize=(12, 6))
sns.lineplot(data=data_original, x="Year", y="National Poverty Rate", label="National Poverty Rate")
sns.lineplot(data=data_original, x="Year", y="Rural Poverty Rate", label="Rural Poverty Rate")
sns.lineplot(data=data_original, x="Year", y="Urban Poverty Rate", label="Urban Poverty Rate")
plt.title("Poverty Rates Over Time")
plt.xlabel("Year")
plt.ylabel("Poverty Rate (%)")
plt.gca().yaxis.set_major_locator(MaxNLocator(integer=True, prune='lower', nbins=5))
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(12, 6))
sns.lineplot(data=data_original, x="Year", y="Literacy Rate", label="Literacy Rate")
plt.title("Literacy Rate Over Time")
```




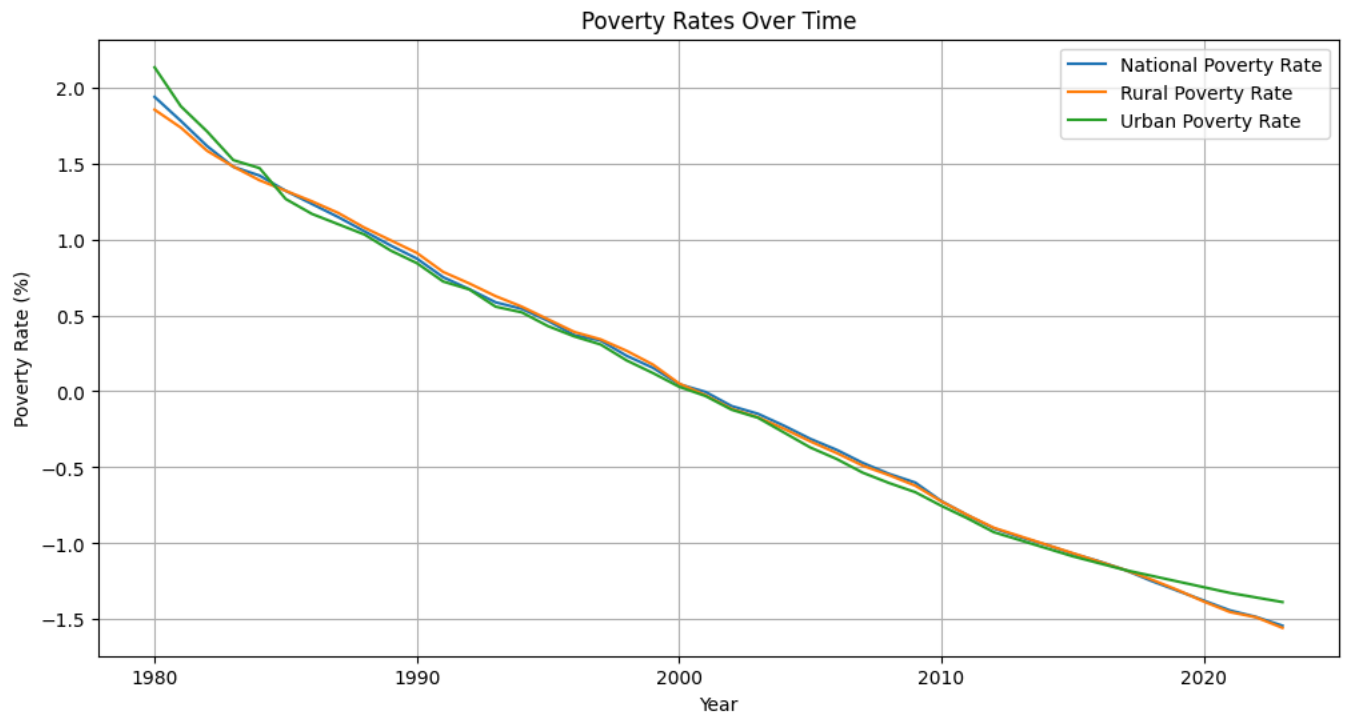
```
plt.xlabel("Year")
plt.ylabel("Literacy Rate (%)")
plt.grid(True)
plt.show()
```




```
# Trend Analysis
plt.figure(figsize=(12, 6))
sns.lineplot(data=data_subset, x="Year", y="National Poverty Rate", label="National Poverty Rate")
sns.lineplot(data=data_subset, x="Year", y="Rural Poverty Rate", label="Rural Poverty Rate")
sns.lineplot(data=data_subset, x="Year", y="Urban Poverty Rate", label="Urban Poverty Rate")
plt.title("Poverty Rates Over Time")
plt.xlabel("Year")
plt.ylabel("Poverty Rate (%)")
plt.legend()
plt.grid(True)
plt.show()


plt.figure(figsize=(12, 6))
```

```
sns.lineplot(data=data_subset, x="Year", y="Literacy Rate", label="Literacy Rate")
plt.title("Literacy Rate Over Time")
plt.xlabel("Year")
plt.ylabel("Literacy Rate (%)")
plt.grid(True)
plt.show()
```




```
# Visualize Rolling Averages
plt.figure(figsize=(12, 6))
sns.lineplot(data=data_subset, x="Year", y="3-Year GDP Rolling Avg", label="3-Year GDP Rolling Avg")
sns.lineplot(data=data_subset, x="Year", y="3-Year Literacy Rolling Avg", label="3-Year Literacy Rolling Avg")
plt.title("3-Year Rolling Averages for GDP and Literacy Rate")
plt.xlabel("Year")
plt.ylabel("Values")
plt.legend()
plt.grid(True)
plt.show()
```






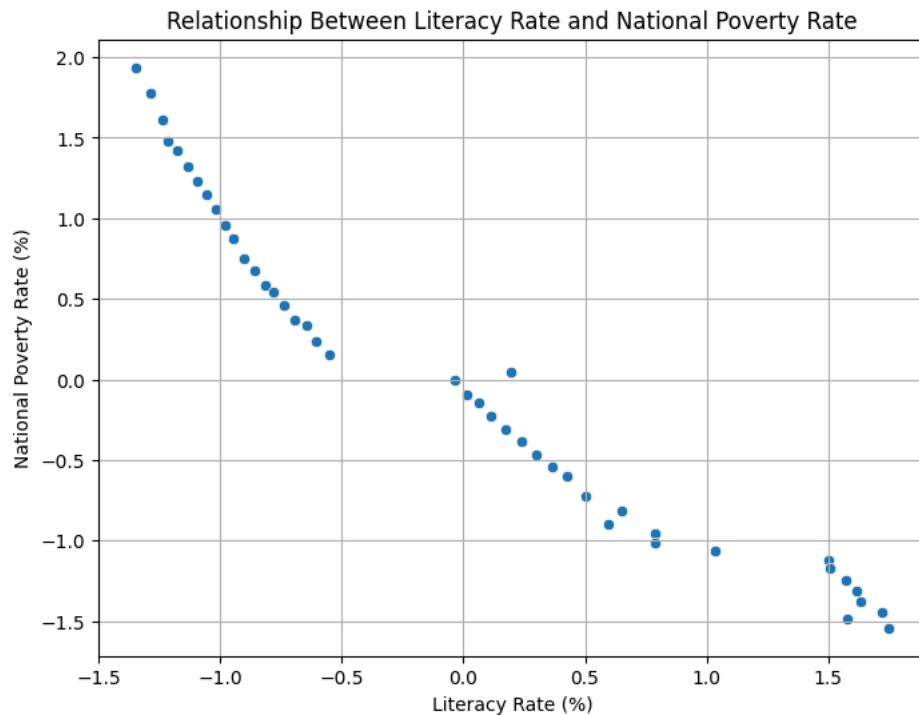
```
# Scatterplot for Interaction Term
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data_subset, x="Urbanization × GDP Growth", y="National Poverty Rate", label="Urbanization × GDP Growth")
plt.title("Impact of Urbanization and GDP Growth on Poverty Rate")
plt.xlabel("Urbanization × GDP Growth")
plt.ylabel("National Poverty Rate (%)")
plt.grid(True)
plt.show()
```



```
# Scatterplots to Analyze Relationships
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data_subset, x="Literacy Rate", y="National Poverty Rate")
plt.title("Relationship Between Literacy Rate and National Poverty Rate")
plt.xlabel("Literacy Rate (%)")
plt.ylabel("National Poverty Rate (%)")
plt.grid(True)
```

```
plt.show()
```

```
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data_subset, x="Education Spending", y="Literacy Rate", label="Education Spending vs Literacy")
sns.scatterplot(data=data_subset, x="Healthcare Spending", y="Literacy Rate", label="Healthcare Spending vs Literacy")
plt.title("Impact of Education and Healthcare Spending on Literacy Rate")
plt.xlabel("Spending (% of GDP / $ Per Capita)")
plt.ylabel("Literacy Rate (%)")
plt.legend()
plt.grid(True)
plt.show()
```




```
# Boxplot for Rolling Literacy Trends vs Poverty Rates
data_subset_melted = data_subset.melt(
    id_vars=["Year", "3-Year Literacy Rolling Avg"],
    value_vars=["Urban Poverty Rate", "Rural Poverty Rate"],
```

```

    var_name="Poverty Type",
    value_name="Poverty Rate"
)
plt.figure(figsize=(10, 6))
sns.boxplot(data=data_subset_melted, x="Poverty Type", y="3-Year Literacy Rolling Avg")
plt.title("3-Year Literacy Rolling Avg vs Urban and Rural Poverty Rates")
plt.xlabel("Poverty Type")
plt.ylabel("3-Year Literacy Rolling Avg")
plt.grid(True)
plt.show()


```



```

# Overlay Trends for Created Features
plt.figure(figsize=(14, 8))
sns.lineplot(data=data_subset, x="Year", y="Urbanization × GDP Growth", label="Urbanization × GDP Growth")
sns.lineplot(data=data_subset, x="Year", y="3-Year GDP Rolling Avg", label="3-Year GDP Rolling Avg")
sns.lineplot(data=data_subset, x="Year", y="3-Year Literacy Rolling Avg", label="3-Year Literacy Rolling Avg")
plt.title("Trends for Created Features Over Time")
plt.xlabel("Year")
plt.ylabel("Values")
plt.legend()
plt.grid(True)
plt.show()

```











```

# create a chart for each feature over time
for col in data_subset.columns[1:]:
    plt.figure(figsize=(5, 4))
    sns.lineplot(data=data_original, x="Year", y=col)
    plt.title(f"{col} Over Time")
    plt.xlabel("Year")
    plt.ylabel(col)
    plt.grid(True)



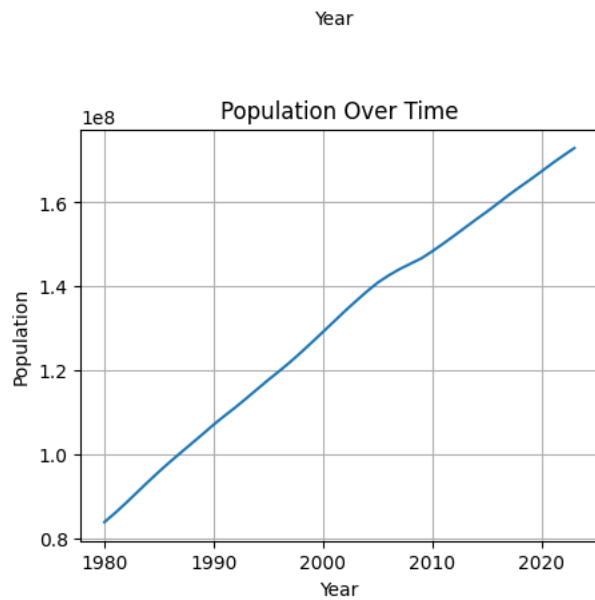
```




```
plt.show()  
print("\n")
```




































### Electricity Access Over Time






```
#scatter plot of literacy rate and unemployment rate
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data_original, x="Literacy Rate", y="Unemployment Rate")
plt.title("Relationship Between Literacy Rate and Unemployment Rate")
plt.xlabel("Literacy Rate")
plt.ylabel("Unemployment Rate")
plt.show()
```



```
#scatter plot of literacy rate and Life expectancy growth rate
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data_original, x="Literacy Rate", y="Life Expectancy Growth Rate")
plt.title("Relationship Between Literacy Rate and Life Expectancy Growth Rate")
plt.xlabel("Literacy Rate")
```




```
plt.figure(figsize=(12, 6))
sns.lineplot(data=data_subset, x="Year", y="Education Spending", label="Education Spending")
sns.lineplot(data=data_subset, x="Year", y="Healthcare Spending", label="Healthcare Spending")
sns.lineplot(data=data_subset, x="Year", y="Infant Mortality Rate", label="Infant Mortality Rate")
```



```

sns.lineplot(data=data_subset, x="Year", y="Maternal Mortality Rate", label="Maternal Mortality Rate")
sns.lineplot(data=data_subset, x="Year", y="Clean Water Access", label="Clean Water Access")
plt.title("Various Socio-economic Measures Over Time")
plt.xlabel("Year")
plt.ylabel("Scaled values")
plt.legend()
plt.show()

```






### National Poverty Rate


National Poverty Rate	1.000000
Rural Poverty Rate	0.999709
Urban Poverty Rate	0.998054
Infant Mortality Rate	0.992928
Death Rate	0.947170
Birth Rate	0.917431
Life Expectancy Growth Rate	0.775513
Net Migration Rate	0.535455
Maternal Mortality Rate	0.302007
Suicide Rate	0.276712
Annual Crime Rate Change	0.071631
Greenhouse Gas Emissions Change	-0.102458
Clean Water Access	-0.275660
Education Spending	-0.301818
Healthcare Spending	-0.311259
Inflation Rate	-0.373114
Urbanization x GDP Growth	-0.396160
Electricity Access	-0.574456
Import Growth	-0.652610
Fossil Fuel Consumption	-0.740503
Export Growth	-0.799517
GDP	-0.814500
GDP Per Capita	-0.818170
3-Year GDP Rolling Avg	-0.819651
Unemployment Rate	-0.907094
Literacy Rate	-0.968302
3-Year Literacy Rolling Avg	-0.973785
Urban Population (%)	-0.979221
Year	-0.997426
Population	-0.998727

**dtype:** float64

```
#scatter plot for National Poverty Rate vs Population
plt.figure(figsize=(10, 6))
sns.scatterplot(data=data_subset, x="Population", y="National Poverty Rate")
plt.title("Relationship Between Population and National Poverty Rate")
plt.xlabel("Population")
plt.ylabel("National Poverty Rate")
plt.show()
```



```
plt.figure(figsize=(12, 6))
sns.lineplot(data=data_subset, x="Year", y="National Poverty Rate", label="National Poverty Rate")
sns.lineplot(data=data_subset, x="Year", y="Literacy Rate", label="Literacy Rate")
sns.lineplot(data=data_subset, x="Year", y="Urban Population (%)", label="Urban Population (%)")
#sns.lineplot(data=data_subset, x="Year", y="Maternal Mortality Rate", label="Maternal Mortality Rate")
#sns.lineplot(data=data_subset, x="Year", y="Clean Water Access", label="Clean Water Access")
plt.title("Socio-economic measures Over Time")
plt.xlabel("Year")
plt.ylabel("Scaled values")
plt.legend()
plt.show()
```



```
# correlation with Literacy Rate
correlation_matrix = data_subset.corr()
correlation_with_literacy = correlation_matrix["Literacy Rate"].sort_values(ascending=False)
print("\nCorrelation with Literacy Rate:")
print(correlation_with_literacy)
```



	1980	1990	2000	2010	2020
Correlation with Literacy Rate:					
Literacy Rate	1.000000				
3-Year Literacy Rolling Avg	0.995405				
Urban Population (%)	0.988159				
Year	0.980226				
Population	0.972824				
Unemployment Rate	0.928044				
GDP Per Capita	0.901279				
GDP	0.897481				
3-Year GDP Rolling Avg	0.888713				
Electricity Access	0.711008				
Export Growth	0.695975				
Fossil Fuel Consumption	0.589086				
Import Growth	0.585368				
Urbanization × GDP Growth	0.535440				
Healthcare Spending	0.416864				
Clean Water Access	0.397836				

```

Inflation Rate      0.294528
Education Spending   0.154902
Greenhouse Gas Emissions Change  0.086434
Annual Crime Rate Change -0.083553
Suicide Rate        -0.305661
Maternal Mortality Rate -0.370832
Net Migration Rate   -0.517146
Death Rate           -0.855011
Life Expectancy Growth Rate -0.861326
Birth Rate           -0.883872
Infant Mortality Rate -0.944369
Urban Poverty Rate   -0.955698
National Poverty Rate -0.968302
Rural Poverty Rate    -0.970977
Name: Literacy Rate, dtype: float64

```

```

# correlation with National Poverty Rate
correlation_matrix = data_subset.corr()
correlation_with_literacy = correlation_matrix["National Poverty Rate"].sort_values(ascending=False)
print("\nCorrelation with National Poverty Rate:")
print(correlation_with_literacy)

```



```

Correlation with National Poverty Rate:
National Poverty Rate      1.000000
Rural Poverty Rate         0.999709
Urban Poverty Rate         0.998054
Infant Mortality Rate      0.992928
Death Rate                 0.947170
Birth Rate                 0.917431
Life Expectancy Growth Rate 0.775513
Net Migration Rate         0.535455
Maternal Mortality Rate    0.302007
Suicide Rate               0.276712
Annual Crime Rate Change   0.071631
Greenhouse Gas Emissions Change -0.102458
Clean Water Access         -0.275660
Education Spending         -0.301818
Healthcare Spending        -0.311259
Inflation Rate             -0.373114
Urbanization × GDP Growth  -0.396160
Electricity Access         -0.574456
Import Growth              -0.652610
Fossil Fuel Consumption    -0.740503
Export Growth              -0.799517
GDP                       -0.814500
GDP Per Capita             -0.818170
3-Year GDP Rolling Avg     -0.819651
Unemployment Rate          -0.907094
Literacy Rate              -0.968302
3-Year Literacy Rolling Avg -0.973785
Urban Population (%)       -0.979221
Year                      -0.997426
Population                 -0.998727
Name: National Poverty Rate, dtype: float64

```

```

# drop features with multicollinearity
drop_multicollinearity = ["GDP", "GDP Per Capita", "Urban Population (%)", "Urban Poverty Rate", "Rural Poverty Rate", "Population", "Fossil Fuel Consumption", "Import Growth", "Death Rate", "Infant Mortality Rate", "Maternal Mortality Rate", "Clean Water Access", "Unemployment Rate", "Birth Rate", "3-Year GDP Rolling Avg", "Literacy Rate"]
data_subset2 = data_subset.drop(columns=drop_multicollinearity)

```

```

# drop features with weak correlation
drop_weak = ["Greenhouse Gas Emissions Change", "Annual Crime Rate Change"]
data_subset3 = data_subset2.drop(columns=drop_weak)

```

```
data_subset3.columns
```



```

Index(['Year', 'National Poverty Rate', 'Education Spending',
       'Healthcare Spending', 'Inflation Rate', 'Net Migration Rate',
       'Export Growth', 'Life Expectancy Growth Rate', 'Electricity Access',
       'Suicide Rate', 'Urbanization × GDP Growth',
       '3-Year Literacy Rolling Avg'],
      dtype='object')


```

```

# Correlation Matrix
correlation_matrix = data_subset3.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=False, cmap="coolwarm", fmt=".2f")

```

```
plt.title("Correlation Matrix (post feature selection)")
plt.show()
```



### 3. Assumption test

```
# missing variable, impute with mean
from sklearn.impute import SimpleImputer

data_subset4 = data_subset3.copy()
imputer = SimpleImputer(strategy="mean")
data_subset4.iloc[:, 1:] = imputer.fit_transform(data_subset3.iloc[:, 1:])

# assumption test - stationarity for each variable
from statsmodels.tsa.stattools import adfuller

for col in data_subset4.columns[1:]:
    print(f"Testing for Stationarity of {col}:")
    result = adfuller(data_subset4[col])
    print(f'ADF Statistic: {result[0]}')
    print(f'p-value: {result[1]}') # p < 0.05 -> Stationary, p > 0.05 -> Not Stationary
```

```
print("\n")
```

```
# create a list of features that are not stationary
non_stationary_features = [col for col in data_subset4.columns[1:] if adfuller(data_subset4[col])[1] > 0.05]
print("Non-Stationary Features:", non_stationary_features)
```

```
Testing for Stationarity of National Poverty Rate:
ADF Statistic: -4.763823203330587
p-value: 6.374826725818937e-05
```

```
Testing for Stationarity of Education Spending:
ADF Statistic: -3.544900041245629
p-value: 0.006906122959191104
```

```
Testing for Stationarity of Healthcare Spending:
ADF Statistic: -1.8008294343679248
p-value: 0.38004998028956016
```

```
Testing for Stationarity of Inflation Rate:
ADF Statistic: -5.145502940672432
p-value: 1.1347447558833673e-05
```

```
Testing for Stationarity of Net Migration Rate:
ADF Statistic: -2.371684077838328
p-value: 0.14988835272038342
```

```
Testing for Stationarity of Export Growth:
ADF Statistic: -1.3197905082476749
p-value: 0.6200702019437832
```

```
Testing for Stationarity of Life Expectancy Growth Rate:
ADF Statistic: -0.6543480750939475
p-value: 0.8582228137788708
```

```
Testing for Stationarity of Electricity Access:
ADF Statistic: -1.4191514377955283
p-value: 0.5730610243262942
```

```
Testing for Stationarity of Suicide Rate:
ADF Statistic: -1.807758642953694
p-value: 0.37661660313197487
```

```
Testing for Stationarity of Urbanization x GDP Growth:
ADF Statistic: -2.276533624733398
p-value: 0.17964806285176094
```

```
Testing for Stationarity of 3-Year Literacy Rolling Avg:
ADF Statistic: 0.28434878043769235
p-value: 0.9766149123955452
```

```
Non-Stationary Features: ['Healthcare Spending', 'Net Migration Rate', 'Export Growth', 'Life Expectancy Growth Rate', 'Elec
```

#### 4. predictive model

```
print(data_subset4.shape)
print(data_subset4.head())
print(data_subset4.info())
```

```
(44, 12)
```

	Year	National Poverty Rate	Education Spending	Healthcare Spending	\
0	1980	1.936390	-2.174833	0.0	
1	1981	1.778304	-1.774808	0.0	
2	1982	1.613032	-1.844885	0.0	
3	1983	1.476503	-1.821526	0.0	
4	1984	1.419017	-1.348503	0.0	
	Inflation Rate	Net Migration Rate	Export Growth	\	
0	-2.175351	-0.404870	-1.244337		
1	0.536530	-0.170065	-1.320995		

```

2      -1.621129      0.063745      -1.331642
3      -0.880755      0.297555      -1.220914
4      -0.491531      0.578128      -1.689379

Life Expectancy Growth Rate  Electricity Access  Suicide Rate \
0      0.252218      0.0      1.634136e-15
1      0.252218      0.0      1.634136e-15
2      0.220236      0.0      1.634136e-15
3      0.188255      0.0      1.634136e-15
4      0.859867      0.0      1.634136e-15

Urbanization × GDP Growth  3-Year Literacy Rolling Avg
0      1.174562      -0.008660
1      1.052558      -0.008660
2      1.028819      -1.290739
3      0.994492      -1.246053
4      0.936896      -1.209492
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 44 entries, 0 to 43
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Year                                44 non-null    int64
1   National Poverty Rate              44 non-null    float64
2   Education Spending                 44 non-null    float64
3   Healthcare Spending                44 non-null    float64
4   Inflation Rate                     44 non-null    float64
5   Net Migration Rate                 44 non-null    float64
6   Export Growth                      44 non-null    float64
7   Life Expectancy Growth Rate        44 non-null    float64
8   Electricity Access                  44 non-null    float64
9   Suicide Rate                       44 non-null    float64
10  Urbanization × GDP Growth           44 non-null    float64
11  3-Year Literacy Rolling Avg         44 non-null    float64
dtypes: float64(11), int64(1)
memory usage: 4.3 KB
None

```

```
#!pip install --upgrade scikit-learn
```

#### ✓ 4.1 Random Forest Model

```

import numpy as np
import pandas as pd
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt

def preprocess_time_series_data(df):
    """Time-series specific preprocessing"""
    # Clean column names and sort by year
    df = df.rename(columns={'Urbanization × GDP Growth': 'Urbanization_GDP_Growth'})
    df = df.sort_values('Year').reset_index(drop=True)

    # Handle zero values in Healthcare Spending
    df['Healthcare Spending'] = np.where(
        (df['Healthcare Spending'] == 0) & (df['Year'] < 2000),
        np.nan,
        df['Healthcare Spending']
    )
    df['Healthcare Spending'] = df['Healthcare Spending'].ffill()

    # Remove non-informative columns
    df = df.drop(columns=['Year', 'Electricity Access', 'Suicide Rate'])

    return df.dropna()

processed_data = preprocess_time_series_data(data_subset4)

# Split features and target
X = processed_data.drop(columns=['National Poverty Rate'])
y = processed_data['National Poverty Rate']

# Time-series aware train-test split (last 20% as test)
test_size = int(len(X) * 0.2)
X_train, X_test = X[:-test_size], X[-test_size:]

```

```

y_train, y_test = y[:-test_size], y[-test_size:]

# Time-series cross-validation configuration
tscv = TimeSeriesSplit(n_splits=5)

# Hyperparameter grid with regularization parameters
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5],
    'max_features': ['sqrt', 0.8],
    'min_impurity_decrease': [0.0, 0.1]
}

# Initialize and tune Random Forest
rf = RandomForestRegressor(random_state=42, n_jobs=-1)
grid_search = GridSearchCV(
    estimator=rf,
    param_grid=param_grid,
    cv=tscv,
    scoring='neg_mean_squared_error',
    verbose=1,
    n_jobs=-1
)

print("Starting grid search with time-series validation...")
grid_search.fit(X_train, y_train)

# Best model evaluation
best_rf = grid_search.best_estimator_

def time_series_evaluate(model, X_train, X_test, y_train, y_test):
    """Time-series specific evaluation"""
    # Generate predictions
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    # Calculate metrics
    metrics = {
        'Train RMSE': np.sqrt(mean_squared_error(y_train, y_pred_train)),
        'Test RMSE': np.sqrt(mean_squared_error(y_test, y_pred_test)),
        'Train MAE': mean_absolute_error(y_train, y_pred_train),
        'Test MAE': mean_absolute_error(y_test, y_pred_test),
        'Train R²': r2_score(y_train, y_pred_train),
        'Test R²': r2_score(y_test, y_pred_test)
    }

    # Plot actual vs predicted
    plt.figure(figsize=(12, 6))
    plt.plot(y_test.values, label='Actual', marker='o')
    plt.plot(y_pred_test, label='Predicted', marker='x')
    plt.title('Time-Series Prediction Performance')
    plt.xlabel('Time Step')
    plt.ylabel('National Poverty Rate')
    plt.legend()
    plt.grid(True)
    plt.show()

    return pd.Series(metrics)

# Evaluate and show results
metrics = time_series_evaluate(best_rf, X_train, X_test, y_train, y_test)
print("\nBest Parameters:", grid_search.best_params_)
print("\nPerformance Metrics:\n", metrics.to_string())


# Feature importance analysis
feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': best_rf.feature_importances_
}).sort_values('Importance', ascending=False)

print("\nFeature Importance:\n", feature_importance.to_string(index=False))

```



Starting grid search with time-series validation...  
Fitting 5 folds for each of 48 candidates, totalling 240 fits



Best Parameters: {'max\_depth': None, 'max\_features': 0.8, 'min\_impurity\_decrease': 0.0, 'min\_samples\_split': 2, 'n\_estimator

#### Performance Metrics:

Train RMSE	0.021702
Test RMSE	0.269561
Train MAE	0.017697
Test MAE	0.245967
Train R <sup>2</sup>	0.997422
Test R <sup>2</sup>	-18.908087

#### Feature Importance:

Feature	Importance
3-Year Literacy Rolling Avg	0.346535
Healthcare Spending	0.320461
Life Expectancy Growth Rate	0.182795
Urbanization_GDP_Growth	0.054611
Net Migration Rate	0.031819
Export Growth	0.023393
Education Spending	0.021718
Inflation Rate	0.018670

## 4.2 Lasso Regression

```
import numpy as np
import pandas as pd
import warnings
from sklearn.model_selection import TimeSeriesSplit, train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.exceptions import ConvergenceWarning
from sklearn.model_selection import GridSearchCV

# Load data and sort by year
data_subset = data_subset4.sort_values('Year')

# Feature Engineering
data_subset['Poverty_Rate_Lag1'] = data_subset['National Poverty Rate'].shift(1)
data_subset['Poverty_Rate_Lag2'] = data_subset['National Poverty Rate'].shift(2)
data_subset['Poverty_Rate_Rolling3'] = (
    data_subset['National Poverty Rate']
    .rolling(window=3, min_periods=1, closed='left') # Ensure no future leakage
    .mean()
)
```

```

# Impute missing values (instead of dropping)
data_subset.fillna(0, inplace=True) # Or use forward-fill

# Define target and predictors
X = data_subset.drop(columns=['National Poverty Rate', 'Year'])
y = data_subset['National Poverty Rate']

# Train-test split (last 20% as test)
test_size = int(0.2 * len(X))
X_train, X_test = X[:-test_size], X[-test_size:]
y_train, y_test = y[:-test_size], y[-test_size:]

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Time-series cross-validation
tscv = TimeSeriesSplit(n_splits=5)

# Hyperparameter tuning with TimeSeriesSplit
param_grid = {'alpha': np.logspace(-5, 3, 50)} # Wider alpha range

with warnings.catch_warnings():
    warnings.filterwarnings("ignore", category=ConvergenceWarning)
    lasso_cv = GridSearchCV(
        Lasso(max_iter=10000),
        param_grid,
        cv=tscv, # Use TimeSeriesSplit
        scoring='neg_mean_squared_error'
    )
    lasso_cv.fit(X_train_scaled, y_train)

# Best model
best_lasso = lasso_cv.best_estimator_

# Predictions and evaluation
y_pred_train = best_lasso.predict(X_train_scaled)
y_pred_test = best_lasso.predict(X_test_scaled)

print("Train Metrics:")
print(f"MAE: {mean_absolute_error(y_train, y_pred_train):.4f}")
print(f"MSE: {mean_squared_error(y_train, y_pred_train):.4f}")
print(f"R²: {r2_score(y_train, y_pred_train):.4f}\n")

print("Test Metrics:")
print(f"MAE: {mean_absolute_error(y_test, y_pred_test):.4f}")
print(f"MSE: {mean_squared_error(y_test, y_pred_test):.4f}")
print(f"R²: {r2_score(y_test, y_pred_test):.4f}")

# Feature importance
feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': best_lasso.coef_
}).sort_values('Coefficient', key=abs, ascending=False)
print("\nFeature Importance:\n", feature_importance)

```

 Train Metrics:  
 MAE: 0.0736  
 MSE: 0.0085  
 R²: 0.9884

Test Metrics:  
 MAE: 3.1801  
 MSE: 16.5092  
 R²: -830.0116

Feature Importance:

	Feature	Coefficient
9	3-Year Literacy Rolling Avg	-1.559530
11	Poverty_Rate_Lag2	-1.258370
10	Poverty_Rate_Lag1	0.962675
12	Poverty_Rate_Rolling3	-0.841784
8	Urbanization × GDP Growth	0.660744
1	Healthcare Spending	-0.426611
3	Net Migration Rate	-0.217992
4	Export Growth	-0.154719
0	Education Spending	-0.129683

```

7           Suicide Rate      -0.090553
5  Life Expectancy Growth Rate  0.082975
6           Electricity Access  0.019748
2           Inflation Rate     -0.004280

```

```
# Residual Plot
```

```
import matplotlib.pyplot as plt
```

```
# Predictions and residuals
```

```
y_pred_test = best_lasso.predict(X_test_scaled)
```

```
residuals = y_test - y_pred_test
```

```
# Plot residuals vs. predicted values
```

```
plt.figure(figsize=(12, 5))
```

```
plt.subplot(1, 2, 1)
```

```
plt.scatter(y_pred_test, residuals, alpha=0.7)
```

```
plt.axhline(0, color='r', linestyle='--')
```

```
plt.title('Residuals vs. Predicted Values')
```

```
plt.xlabel('Predicted Poverty Rate')
```

```
plt.ylabel('Residuals')
```

```
# Plot residuals over time (test set years)
```

```
test_years = data_subset['Year'].iloc[-len(y_test):]
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(test_years, residuals, marker='o', linestyle='-', color='b')
```

```
plt.axhline(0, color='r', linestyle='--')
```

```
plt.title('Residuals Over Time')
```



```
plt.xlabel('Year')
```

```
plt.ylabel('Residuals')
```

```
plt.xticks(rotation=45)
```

```
plt.tight_layout()
```

```
plt.show()
```



### 4.3 VAR Model

```

import numpy as np
import pandas as pd
from statsmodels.tsa.api import VAR
from statsmodels.tsa.stattools import adfuller, grangercausalitytests
from statsmodels.stats.stattools import durbin_watson
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt

```

```
# =====
```

```
# 1. Data Preparation & Cleaning
```

```
# =====
```

```
data = data_subset4.sort_values('Year').set_index('Year')
```

```
data.replace(0, np.nan, inplace=True)
```

```

data.fillna(method='ffill', inplace=True)
data = data.drop(columns=['Electricity Access', 'Suicide Rate'])

# =====
# 2. Stationarity Transformation (Dynamic Handling)
# =====
def make_stationary(df, target_col):
    stationary_df = df.copy()
    # Check if target needs differencing
    p_value_target = adfuller(df[target_col].dropna())[1]
    if p_value_target > 0.05:
        stationary_df[f"{target_col}_diff"] = df[target_col].diff()
        stationary_df.drop(target_col, axis=1, inplace=True)
        target_var = f"{target_col}_diff"
    else:
        target_var = target_col

    # Check other features
    for col in df.columns:
        if col != target_col:
            p_value = adfuller(df[col].dropna())[1]
            if p_value > 0.05:
                stationary_df[f"{col}_diff"] = df[col].diff()
                stationary_df.drop(col, axis=1, inplace=True)
    return stationary_df.dropna(), target_var

stationary_data, target_var = make_stationary(data, 'National Poverty Rate')

# =====
# 3. Feature Selection (Granger Causality) - Updated
# =====
candidate_features = [col for col in stationary_data.columns if col != target_var]
significant_features = []

print("\nGranger Causality Results:")
for feature in candidate_features:
    test_result = grangercausalitytests(stationary_data[[target_var, feature]], maxlag=2, verbose=False)
    p_values = [test_result[i+1][0]['ssr_chi2test'][1] for i in range(2)]
    min_p = min(p_values)
    print(f"{feature}: min p-value = {min_p:.4f}")
    if min_p < 0.1: # Relaxed threshold
        significant_features.append(feature)

final_features = [target_var] + significant_features
filtered_data = stationary_data[final_features]

# Add variance filtering
from sklearn.feature_selection import VarianceThreshold
selector = VarianceThreshold(threshold=0.01)
filtered_data = pd.DataFrame(
    selector.fit_transform(filtered_data),
    columns=filtered_data.columns[selector.get_support()]
)

print("\nFinal Features After Filtering:", filtered_data.columns.tolist())

# =====
# 4. Train-Test Split & Model Fitting
# =====
test_size = 5
train = filtered_data.iloc[:-test_size]
test = filtered_data.iloc[-test_size:]

model = VAR(train)
lag_order = model.select_order(maxlags=3).selected_orders['aic']
var_model = model.fit(lag_order)

# =====
# 5. Feature Importance Analysis (Updated)
# =====
def calculate_var_feature_importance(model, target_index):
    """Calculate and format feature importance with full feature listing"""
    coeffs = model.coefs
    importance = {feature: 0.0 for feature in train.columns if feature != target_var}

    # Calculate raw importance scores
    for lag in range(lag_order):

```